

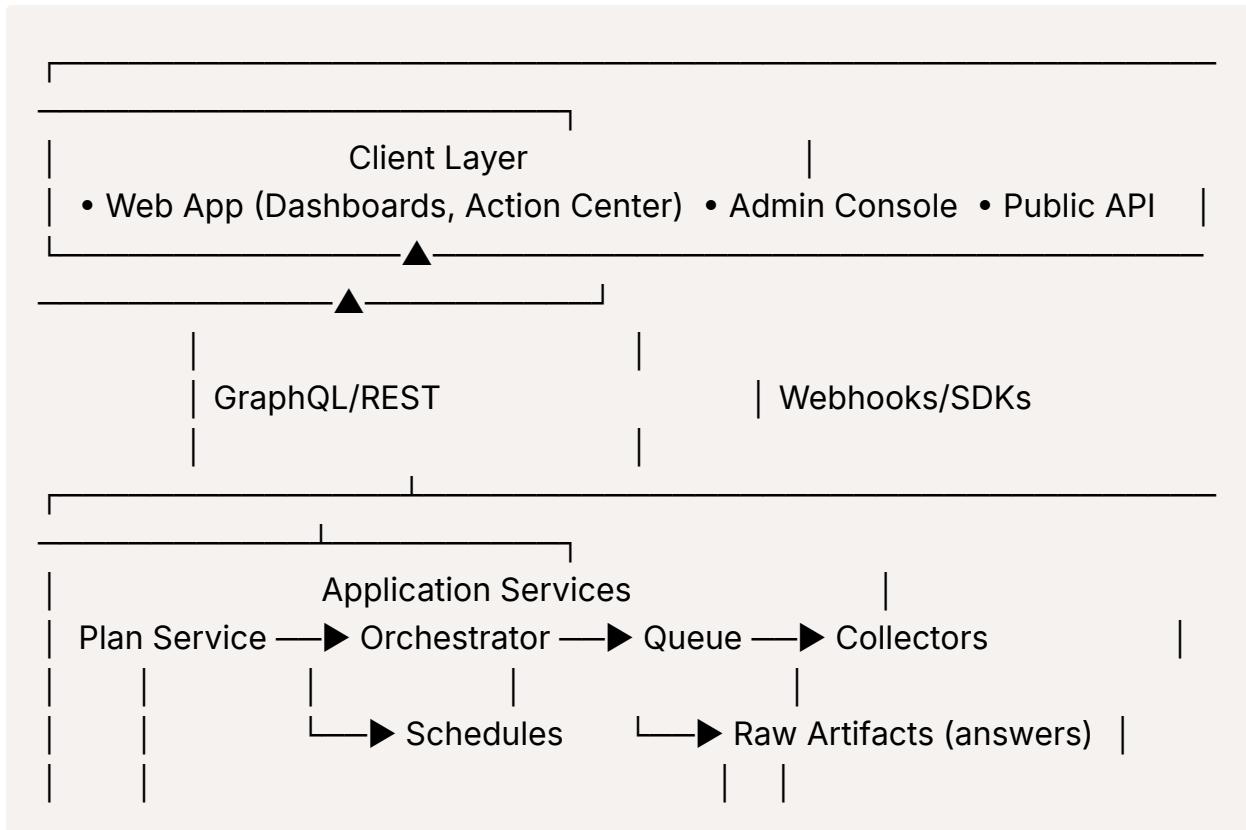
# GEO/AI Search Monitoring & Optimization Platform

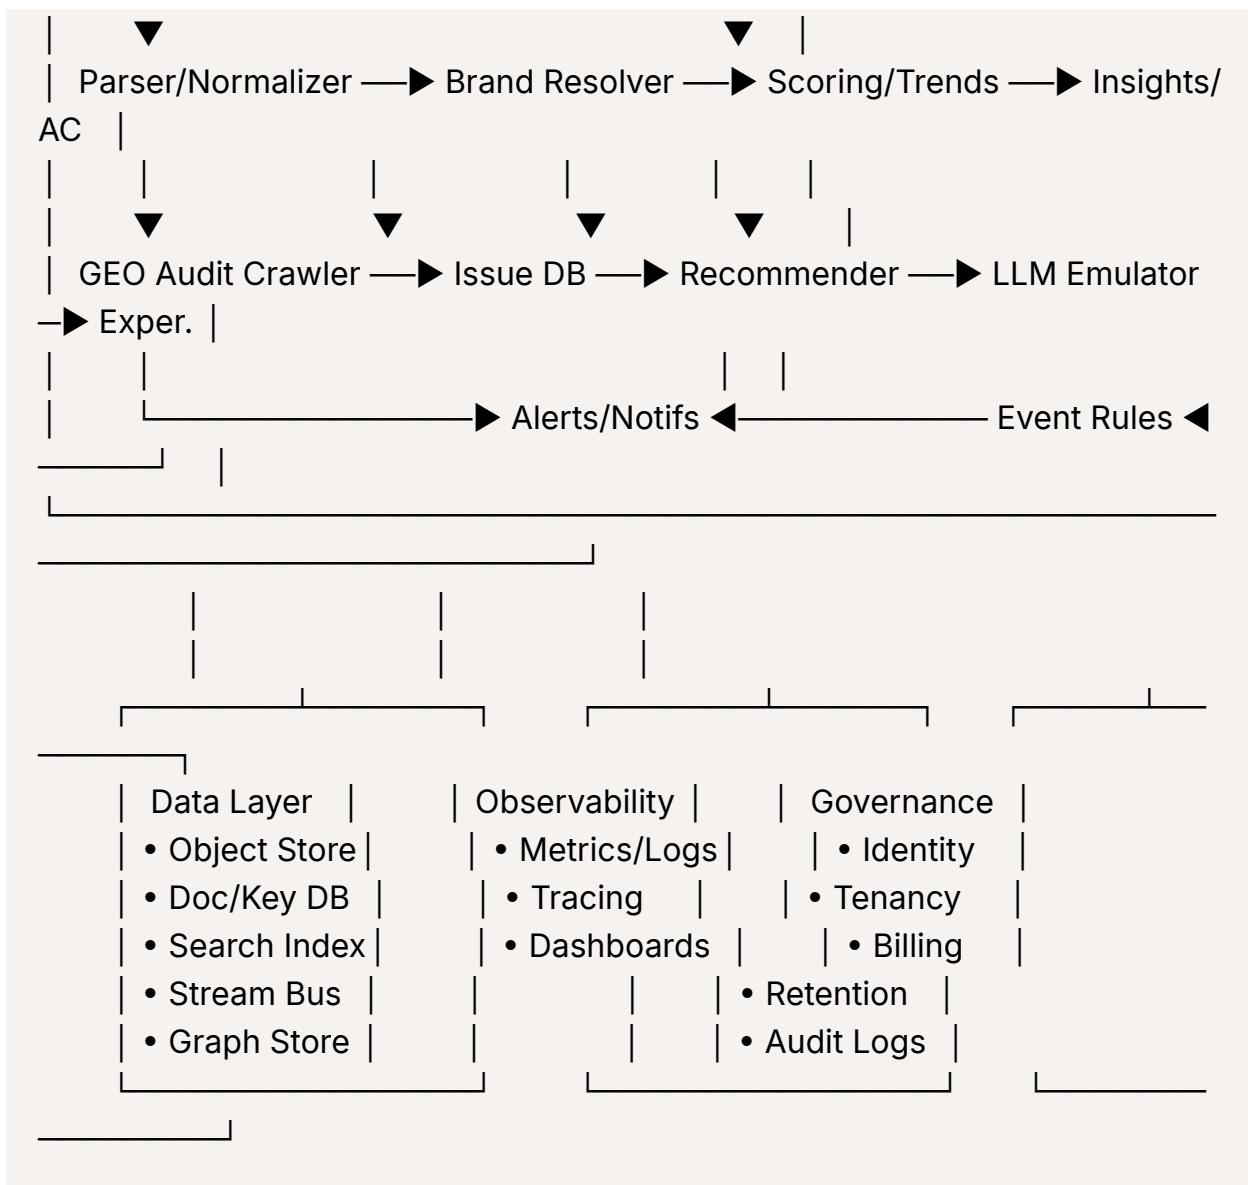
Created by	Bibo Wang
Created time	@October 24, 2025 12:04 PM
Tags	

## Cloud-Agnostic System Architecture (Canvas)

Purpose: A vendor-neutral blueprint tying business capabilities to technical components. Replace terms like "object storage" or "stream bus" with your preferred tech later.

### 1) High-Level Architecture (Conceptual)





**Legend:** AC = Action Center; Exper. = Experiment Lab.

## 2) Component Catalog (Business ↔ Tech)

### Client Layer

- **Web App (Dashboards & Workbench):** Explore visibility, evidence (answers + citations), competitor diffs; manage prompts, actions, experiments.
- **Admin Console:** Workspaces, roles, usage, plans, retention.

- **Public API/SDK:** Read scores/answers/audits; push prompts; receive webhooks.

## Core Application Services

### 1. Plan Service (Discover & Plan)

*Business:* Define topics, prompts, engines, locales, cadence, budget.

*Tech:* REST/GraphQL for CRUD; validation; versioned plan snapshots.

### 2. Orchestrator (Run Scheduling)

*Business:* Turns plans into executable jobs.

*Tech:* Scheduler → Job generator → enqueues to **Queue** with back-pressure.

### 3. Queue (Work Dispatch)

*Business:* Smooths spikes; isolates failures.

*Tech:* Durable FIFO/standard messages, retry & dead-letter.

### 4. Collectors (Acquisition)

*Business:* Ask AI engines; capture answers, citations, screenshots.

*Tech:* Headless browser/HTTP clients, proxy rotation, human-like pacing; write **Raw Artifacts** to **Object Storage**.

### 5. Parser & Normalizer

*Business:* Make results comparable across engines.

*Tech:* Async workers transform HTML/JSON → `AnswerRecord` + `Citation[]`.

### 6. Brand & Competitor Resolver

*Business:* Map mentions/links to brands & competitors.

*Tech:* Alias dictionaries, fuzzy matching, ownership graph.

### 7. Scoring & Trends

*Business:* KPIs: Answer Share, Citation Share, Prominence, Sentiment, Deltas.

*Tech:* Batch jobs computing daily `ScoreSnapshot`; indexed in **Search**.

### 8. Insights & Action Center

*Business:* Prioritized fixes with expected impact (content outlines, schema tasks).

*Tech:* Heuristics + simple models; LLM assist for outlines; ticketing/webhook connectors.

## 9. GEO Audit Crawler

*Business:* Site readiness check (FAQ/schema/metadata/speed/clarity).

*Tech:* Rendered crawl; rules engine emits `AuditIssue` + fix recipes.

## 10. LLM Emulator (Test Harness)

*Business:* Predict which variant is more cite-worthy before publishing.

*Tech:* Deterministic LLM configs, judge prompts, reproducible seeds.

## 11. Experiment Lab

*Business:* Validate lift with pre/post or A/B; publish proof.

*Tech:* Experiment registry, windows, significance/uplift calc.

## 12. Alerts & Notifications

*Business:* Digest + real-time alerts on drops/spikes & competitor surges.

*Tech:* Rule engine on metric streams → email/Chat/webhooks.

# Data Layer (Technology-Agnostic)

- **Object Storage:** Raw HTML/JSON/PNG artifacts; versioned; lifecycle policies.
- **Document/Key-Value DB:** Jobs, answers, citations, scores, audits, actions, experiments.
- **Search Index:** Fast filtering/aggregations (by engine/locale/topic/brand/date).
- **Stream Bus:** Event pipeline (ingest → parse → resolve → score → index).
- **Graph Store (optional):** Brand ↔ Domain ↔ Topic ↔ Citation relationships.
- **Cache:** Hot queries, pre-computed leaderboards.

# Observability & Governance

- **Metrics/Logging/Tracing:** p95 latencies, run success, parse errors, cost/run.

- **Identity & Tenancy:** Workspaces, roles (Viewer/Editor/Admin), API tokens.
  - **Billing/Usage:** Runs, seats, emulator hours; plan limits, overage.
  - **Data Retention & Audit:** Configurable TTLs; immutable audit logs.
- 

### 3) Key Data Contracts (neutral JSON)

#### QueryPlan

```
{  
  "plan_id": "qp_2025_10_24",  
  "workspace_id": "ws_...",  
  "engines": ["chatgpt","perplexity","gemini"],  
  "locales": ["en-US"],  
  "prompt_ids": ["p1","p2"],  
  "cadence": "daily",  
  "budget_cap": {"runs_per_day": 800}  
}
```

#### AnswerRecord

```
{  
  "answer_id": "ans_...",  
  "job_id": "job_...",  
  "engine": "chatgpt",  
  "prompt_id": "p1",  
  "locale": "en-US",  
  "answer_text": "markdown...",  
  "answer_type": "summary",  
  "citations": [  
    {"url": "https://brand.com/x", "domain": "brand.com", "position": 1, "first_party": true}  
  ],  
  "evidence": [{"type": "screenshot", "uri": "obj://raw/.../page.png"}],  
}
```

```
        "created_at": "2025-10-24T09:00:00Z"  
    }  

```

## ScoreSnapshot (daily)

```
{  
    "score_id": "sc_...",  
    "workspace_id": "ws_...",  
    "brand_id": "br_...",  
    "topic_id": "t_...",  
    "engine": "chatgpt",  
    "locale": "en-US",  
    "window_start": "2025-10-23",  
    "answer_share": 0.42,  
    "citation_share": 0.35,  
    "prominence": 0.61,  
    "sentiment_avg": 0.18,  
    "delta_answer_share": 0.07  
}  

```

## AuditIssue

```
{  
    "audit_id": "au_...",  
    "brand_id": "br_...",  
    "url": "https://brand.com/x",  
    "issue_code": "FAQ_SCHEMA_MISSING",  
    "severity": "high",  
    "rationale": "FAQ block absent; competitors cited for Q&A queries",  
    "fix_recipe": {"type": "faq_block", "schema": "FAQPage", "example": "..."},  
    "impacted_prompts": ["p1", "p7"],  
    "status": "open"  
}  

```

## 4) Core Workflows (Step-by-Step, Plain Language)

### A) Daily Monitoring Loop

1. **Plans** create jobs → **Orchestrator** schedules them.
2. **Collectors** ask AI engines; store raw pages & screenshots.
3. **Parser** extracts answer text + links; **Resolver** tags brands.
4. **Scoring** summarizes KPIs; **Index** makes them searchable.
5. **Dashboard** shows trends; **Alerts** flag anomalies.

### B) GEO Audit → Action Center

1. **Crawler** scans selected site paths; runs checks.
2. Emits **AuditIssues** with severity + how-to fix.
3. **Recommender** ties issues to lost prompts & expected lift.
4. One-click create tasks or CMS drafts via **Integrations**.

### C) LLM Emulator (Pre-Publish Test)

1. Choose a losing prompt + page; load Variant A/B.
2. **Emulator** asks a judge model which is more cite-worthy.
3. Store result; ship the predicted winner.

### D) Experiment Lab (Post-Publish Validation)

1. Define window & pages; run pre/post or A/B.
2. Collect metrics; compute **uplift** with significance.
3. Publish report with evidence.

### E) Communications

- **Digest:** weekly summary of wins/losses, actions completed.
- **Alerts:** immediate notifications for drops/surges.
- **Exports/API:** feed BI or CRM with scores & issues.

---

## 5) Non-Functional Requirements (Targets)

- **Reliability:**  $\geq 99.9\%$  dashboard availability;  $\geq 97\%$  successful runs/day.
  - **Performance:** p95 query < 1.5 s; ingestion to dashboard < 60 min.
  - **Cost Controls:** per-workspace budgets; back-pressure on queue.
  - **Security:** org/workspace isolation; encrypted data at rest & transit; signed evidence URLs.
  - **Privacy/Robots:** configurable modes (strict/standard/lab) with customer consent.
- 

## 6) Deployment Topology (Technology-Neutral)

- **Stateless Services:** Orchestrator, Parsers, Scorers, Recommender, API.
  - **Worker Pools:** Collectors, Crawler, Emulator jobs (container or function runtime).
  - **Data Stores:** Object storage, document DB, search index, stream bus, optional graph store.
  - **Gateways:** API (REST/GraphQL), Webhook receiver, Auth server.
  - **Pipelines:** CI/CD with automated tests & canary releases.
- 

## 7) Roadmap (Business + Build)

- **Phase 1 (MVP):** Plans, Collectors (2 engines), Parser, Dashboard v1, GEO Audit v1, Alerts.
  - **Phase 2:** Brand Resolver, Scoring/Trends, Action Center, API/Exports, Dashboard v2.
  - **Phase 3:** LLM Emulator, Experiment Lab, Graph Explorer.
  - **Phase 4:** Automation rules, Billing/Tenancy, Compliance toolkit.
- 

## 8) Glossary

- **Answer Share:** % prompts where your brand appears in AI answers.
  - **Citation Share:** % of links in AI answers pointing to your domains.
  - **Prominence:** Weighting of mention order/position in the answer.
  - **GEO:** Generative Engine Optimization — optimizing to be cited by AI.
- 

## Notes

- Replace each generic component with your chosen technology (e.g., "object storage" → your blob store).
- This canvas is intentionally implementation-agnostic so it fits any cloud/on-prem stack.